

Probabilistic Neural Architecture Search

Chen Zhao

(Arxiv 2019, <https://arxiv.org/pdf/1902.05116.pdf>)

1 Motivations and Contributions

- The trade-off between high computational cost and high memory cost has resulted in the usage of different types of surrogates (or proxies) during architecture search, architecture surrogates and dataset surrogates.
 - Architecture surrogates are small versions of the final network that are cheaper to store in memory and faster to train.
 - Dataset surrogates are datasets that can be used as a proxy to perform a quicker search. For example, one could perform architecture search on CIFAR-10 and "transfer" the resulting architecture (with some modifications, e.g. changing the number of filters) to ImageNet.
- While good performance on a surrogate task does not guarantee good performance on the final task (i.e. fully-sized network on the target dataset), the vast majority of architecture search methods use at least one, if not both, of these surrogates.
- The proposed procedure PARSEC (Probabilistic neural ARchitecture SEarCh) is memory-efficient.
- PARSEC transfers probability distributions over architectures learnt on small surrogates to larger networks and datasets, enabling us to further reduce the computational cost.
- PARSEC is computationally efficient and can be run in less than a day on a single GPU on CIFAR-10.
- PARSEC outperforms other methods that consider the same architecture search space, while drastically reducing the search time.

2 Architecture Search Space

In this work, we consider the same architecture search space as in DARTS. Neural network architectures are obtained by stacking a recurrent convolutional unit, denoted as cell. A cell is defined as a directed acyclic graph of N ordered

nodes, $\{\mathbf{z}_1, \dots, \mathbf{z}_N\}$. An intermediate node, \mathbf{z}_k , is defined as the sum of two of the previous nodes, \mathbf{z}_i and \mathbf{z}_j with $i, j < k$, after applying primitive operations $o_{i,k}$ and $o_{j,k}$, respectively:

$$\mathbf{z}_k = o_{i,k}(\mathbf{z}_i) + o_{j,k}(\mathbf{z}_j) \quad (1)$$

3 Deterministic Approach to Architecture Search

It consider an over-parametrized parent network containing all possible paths between nodes:

$$\mathbf{z}_k = \sum_{i,p} \alpha_{i,k}^p \cdot o_{i,k}^p(\mathbf{z}_i) \quad \text{where} \quad \sum_p \alpha_{i,k}^p = 1 \quad (2)$$

where indices i and p run over all possible inputs and operations for node k respectively. $o_{i,k}^p$ denotes primitive operation p on input node i to contribute to output node k , and $\alpha_{i,k}^p$ is the weight of the path.

4 Probabilistic Neural Architecture Search

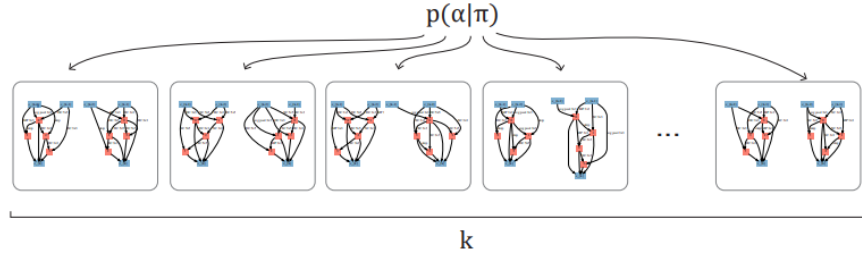


Figure 1: Pictorial representation of architecture sampling. Each sample consists of a normal and a reduction cell, which are shown side-to-side.

In PARSEC, child networks from the search space are directly obtained as samples from the specified architecture distribution, and thus neural architecture is reduced to inferring the probability distribution $p(\alpha|\pi)$ of high-performing architectures on the task at hand. Given a supervised task and denoting with \mathbf{y} the targets and with \mathbf{X} the input features, this is achieved by optimizing the continuous prior hyper-parameters π through an empirical Bayes Monte Carlo procedure. We optimize

$$p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \pi) = \int p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \alpha) p(\alpha|\pi) d\alpha \quad (3)$$

with respect to the network weights \mathbf{v} , which are shared across all child architectures (i.e., samples), and the prior hyper-parameters π , i.e., the architecture parameters.

We assume independent categorical distributions over each input/operation pair. Specifically, for normal and reduction cells we set

$$\begin{aligned} p(\boldsymbol{\alpha}^{(n)}|\boldsymbol{\pi}^{(n)}) &= \prod_{n=1}^N \prod_{i=1}^2 \text{Cat}(\boldsymbol{\alpha}_{n,i}^{(n)}|\boldsymbol{\pi}_{n,i}^{(n)}), \\ p(\boldsymbol{\alpha}^{(r)}|\boldsymbol{\pi}^{(r)}) &= \prod_{n=1}^N \prod_{i=1}^2 \text{Cat}(\boldsymbol{\alpha}_{n,i}^{(r)}|\boldsymbol{\pi}_{n,i}^{(r)}), \end{aligned}$$

where n runs over the nodes, i runs over the inputs of each node (in our search space each node has two inputs), $\boldsymbol{\pi}_{n,i}^{(n)}$ and $\boldsymbol{\pi}_{n,i}^{(r)}$ are the vectors of probabilities for all possible incoming node/operation pairs and finally, we introduced $\boldsymbol{\alpha}^{(\cdot)} = \{\boldsymbol{\alpha}_{n,i}^{(\cdot)}\}$ and $\boldsymbol{\pi}^{(\cdot)} = \{\boldsymbol{\pi}_{n,i}^{(\cdot)}\}$.

5 Importance-Weighted Monte Carlo empirical Bayes

We develop an importance weighted EB procedure for jointly optimizing π and v . We begin by introducing the following estimator:

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) &= \int p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\alpha}) p(\boldsymbol{\alpha}|\boldsymbol{\pi}) d\boldsymbol{\alpha} \\ &\approx \frac{1}{K} \sum_k p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\alpha}_k), \end{aligned}$$

with $\boldsymbol{\alpha}_k \sim p(\boldsymbol{\alpha}|\boldsymbol{\pi})$. Note that the gradients can be written as

$$\begin{aligned} \nabla_{\mathbf{v}, \boldsymbol{\pi}} \log p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) &= \frac{1}{p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi})} \int \nabla_{\mathbf{v}, \boldsymbol{\pi}} p(\mathbf{y}, \boldsymbol{\alpha}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) d\boldsymbol{\alpha} \\ &= \frac{1}{p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi})} \int \nabla_{\mathbf{v}, \boldsymbol{\pi}} \log p(\mathbf{y}, \boldsymbol{\alpha}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) p(\mathbf{y}, \boldsymbol{\alpha}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) d\boldsymbol{\alpha} \end{aligned}$$

Finally, taking an importance sampling estimator to the expectation, we obtain tractable gradient estimators for θ and π :

$$\begin{aligned}\nabla_{\mathbf{v}} \log p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) &\approx \sum_{k=1}^K \tilde{\omega}_k \nabla_{\mathbf{v}} \log p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\alpha}_k) \triangleq \tilde{\nabla}_{\mathbf{v}}, \\ \nabla_{\boldsymbol{\pi}} \log p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \boldsymbol{\pi}) &\approx \sum_{k=1}^K \tilde{\omega}_k \nabla_{\boldsymbol{\pi}} \log p(\boldsymbol{\alpha}_k|\boldsymbol{\pi}) \triangleq \tilde{\nabla}_{\boldsymbol{\pi}},\end{aligned}$$

where $\alpha_k \sim p(\alpha|\pi)$ and $\tilde{\omega}_k = \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \alpha_k)}{\sum_j p(\mathbf{y}|\mathbf{X}, \mathbf{v}, \alpha_j)}$.

Algorithm 1 Importance weighted Monte-Carlo EB algorithm used for joint training of the network and architecture parameters.

```

1: Define train and search splits of the data
2:  $\boldsymbol{\theta}, \boldsymbol{\pi} \leftarrow$  Initial values
3: Initialize  $\boldsymbol{\theta}$  and  $\boldsymbol{\pi}$ 
4: repeat
5:   Sample batch  $(\mathbf{X}^{(s)}, \mathbf{y}^{(s)})$  from search set
6:   for  $k \in \{1, \dots, K\}$  do
7:     Sample architecture  $\boldsymbol{\alpha}_k$  from  $p(\boldsymbol{\alpha}|\boldsymbol{\pi})$ 
8:      $\omega_k \leftarrow p(\mathbf{y}^{(s)}|\mathbf{X}^{(s)}, \mathbf{v}, \boldsymbol{\alpha}_k)$ 
9:   end for
10:  Compute normalized weights  $\tilde{\omega}_k = \frac{\omega_k}{\sum_{j=1}^K \omega_j}$ 
11:  Sample batch  $(\mathbf{X}^{(t)}, \mathbf{y}^{(t)})$  from train set
12:  for  $k \in \{1, \dots, K\}$  do
13:     $g_{\mathbf{v},k} \leftarrow \nabla_{\mathbf{v}} \log p(\mathbf{y}^{(t)}|\mathbf{X}^{(t)}, \mathbf{v}, \boldsymbol{\alpha}_k)$ 
14:     $g_{\boldsymbol{\pi},k} \leftarrow \nabla_{\boldsymbol{\pi}} \log p(\boldsymbol{\alpha}_k|\boldsymbol{\pi})$ 
15:  end for
16:  Update  $\mathbf{v}$  by  $\tilde{\nabla}_{\mathbf{v}} = \sum_k \tilde{\omega}_k g_{\mathbf{v},k}$ 
17:  Update  $\boldsymbol{\pi}$  by  $\tilde{\nabla}_{\boldsymbol{\pi}} = \sum_k \tilde{\omega}_k g_{\boldsymbol{\pi},k}$ 
18: until Number of epochs is reached or convergence is achieved
19: Return:  $(\boldsymbol{\theta}, \boldsymbol{\pi})$ 

```
